

Cost-Aware Vertical Scaling Under Varying Workload

Hardware Configuration

low load

medium load

high load

peak load

sustain peak

ramp-down 1

ramp-down 2

ramp-down 3

low load

cpu_4 cpu_12 gpu_25 gpu_100

gpu_100

gpu_25

cpu_12

cpu_4

Time (minutes)

Throughput EMA (tok/s)

4000

3000

2000

1000

0

0

20

40

60

80

100

120

147

52